



MC²: High Performance GC for Memory-Constrained Environments

Narendran Sachindran

J. Eliot B. Moss

Emery D. Berger

University of Massachusetts Amherst



Motivation

- Java widely used
 - Safety
 - Portability
- Garbage collector requirements
 - High throughput
 - Short pauses
 - Good memory utilization



Motivation

- Handheld Devices
 - Cellular phones, PDAs widely used
 - Constrained memory
- Diverse applications
 - Media players
 - Video games
 - Digital cameras
 - GPS
 - Scaled down desktop apps (e-mail, browser etc.)
- Require high throughput, short response time



Talk Outline

- Generational collection
- MC² overview
- Algorithmic Details
- Experimental Results
- Conclusions




Generational Collection

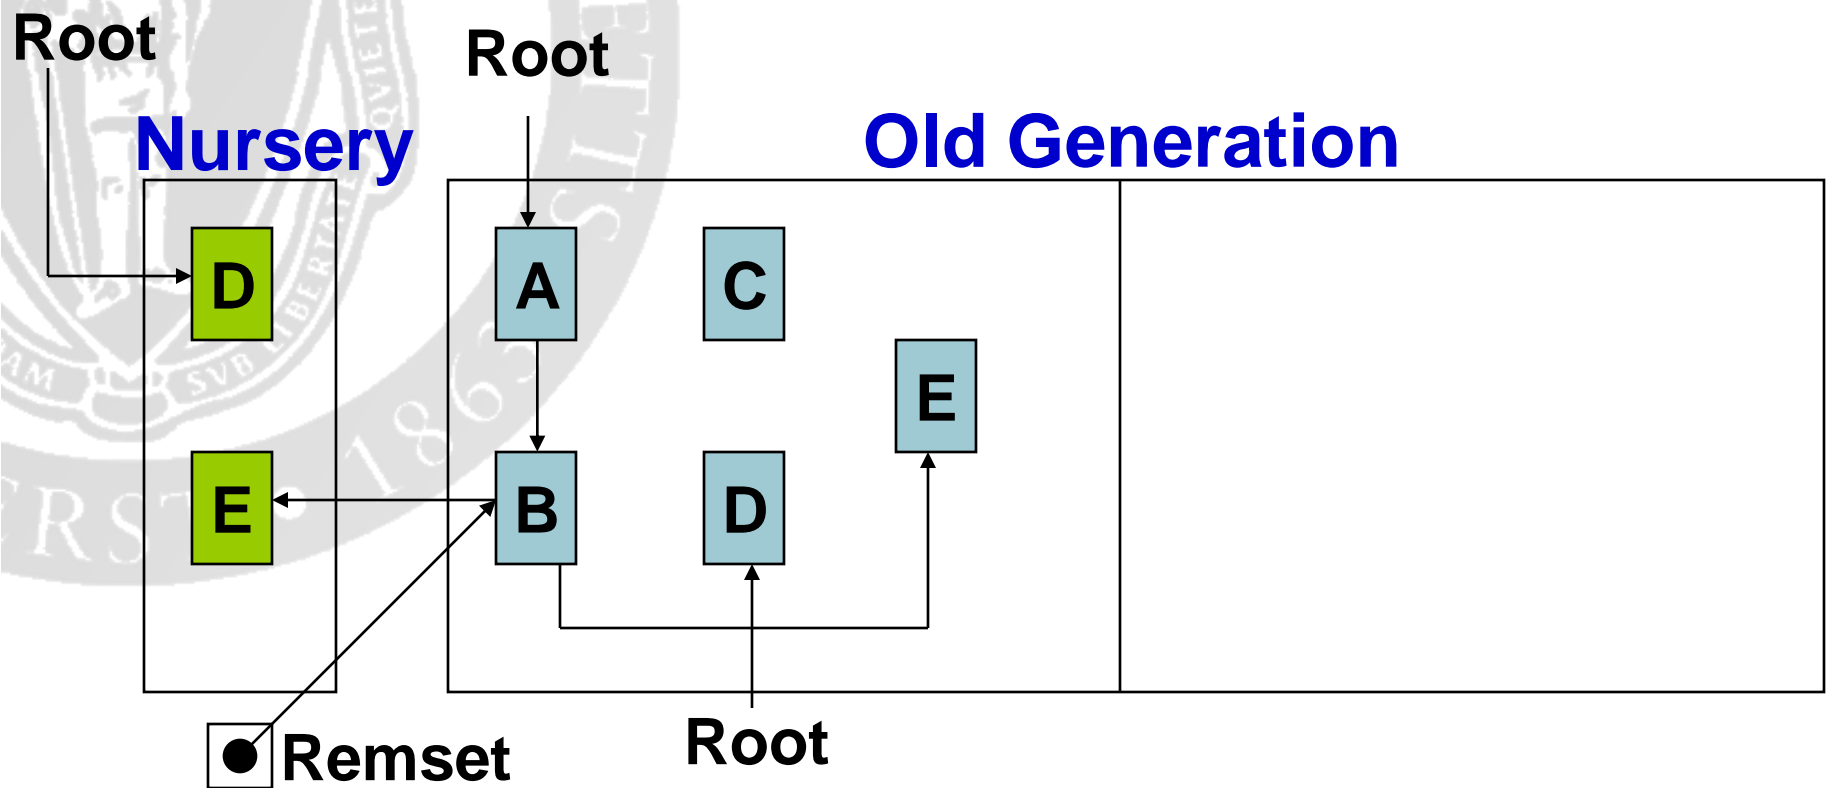
- Divide heap into regions called generations
- Generations segregate objects by age
- Focus GC effort on younger objects

Nursery

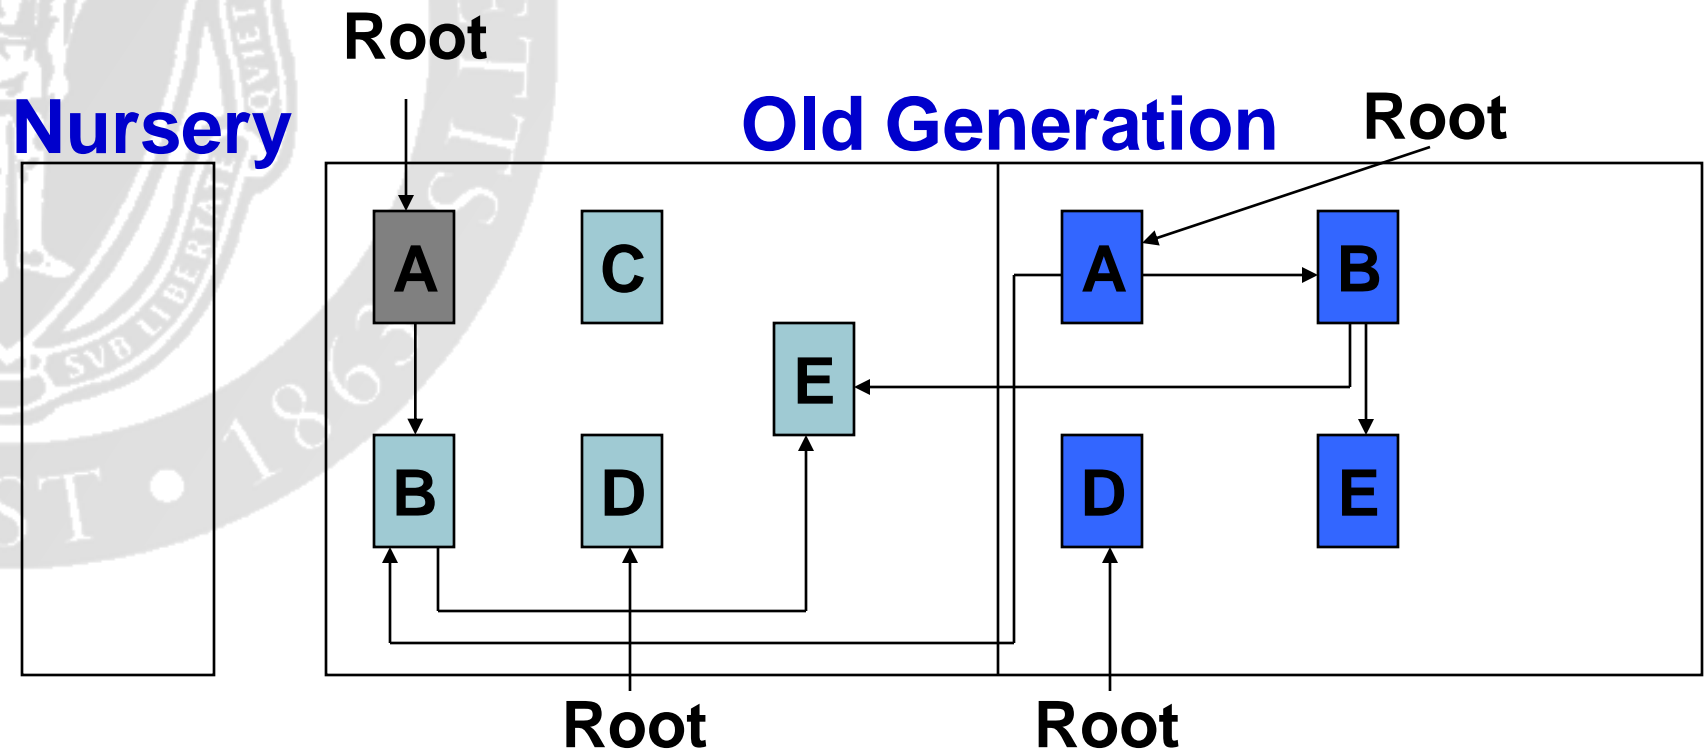
Old Generation



Generational Copying Collection




Generational Copying Collection



MC² Overview

- Extends gen. copying, overcomes 2X overhead
- Divides space into equal size *windows*
- Reserves one or more windows for copying
- Collects in two phases: *mark* and *copy*

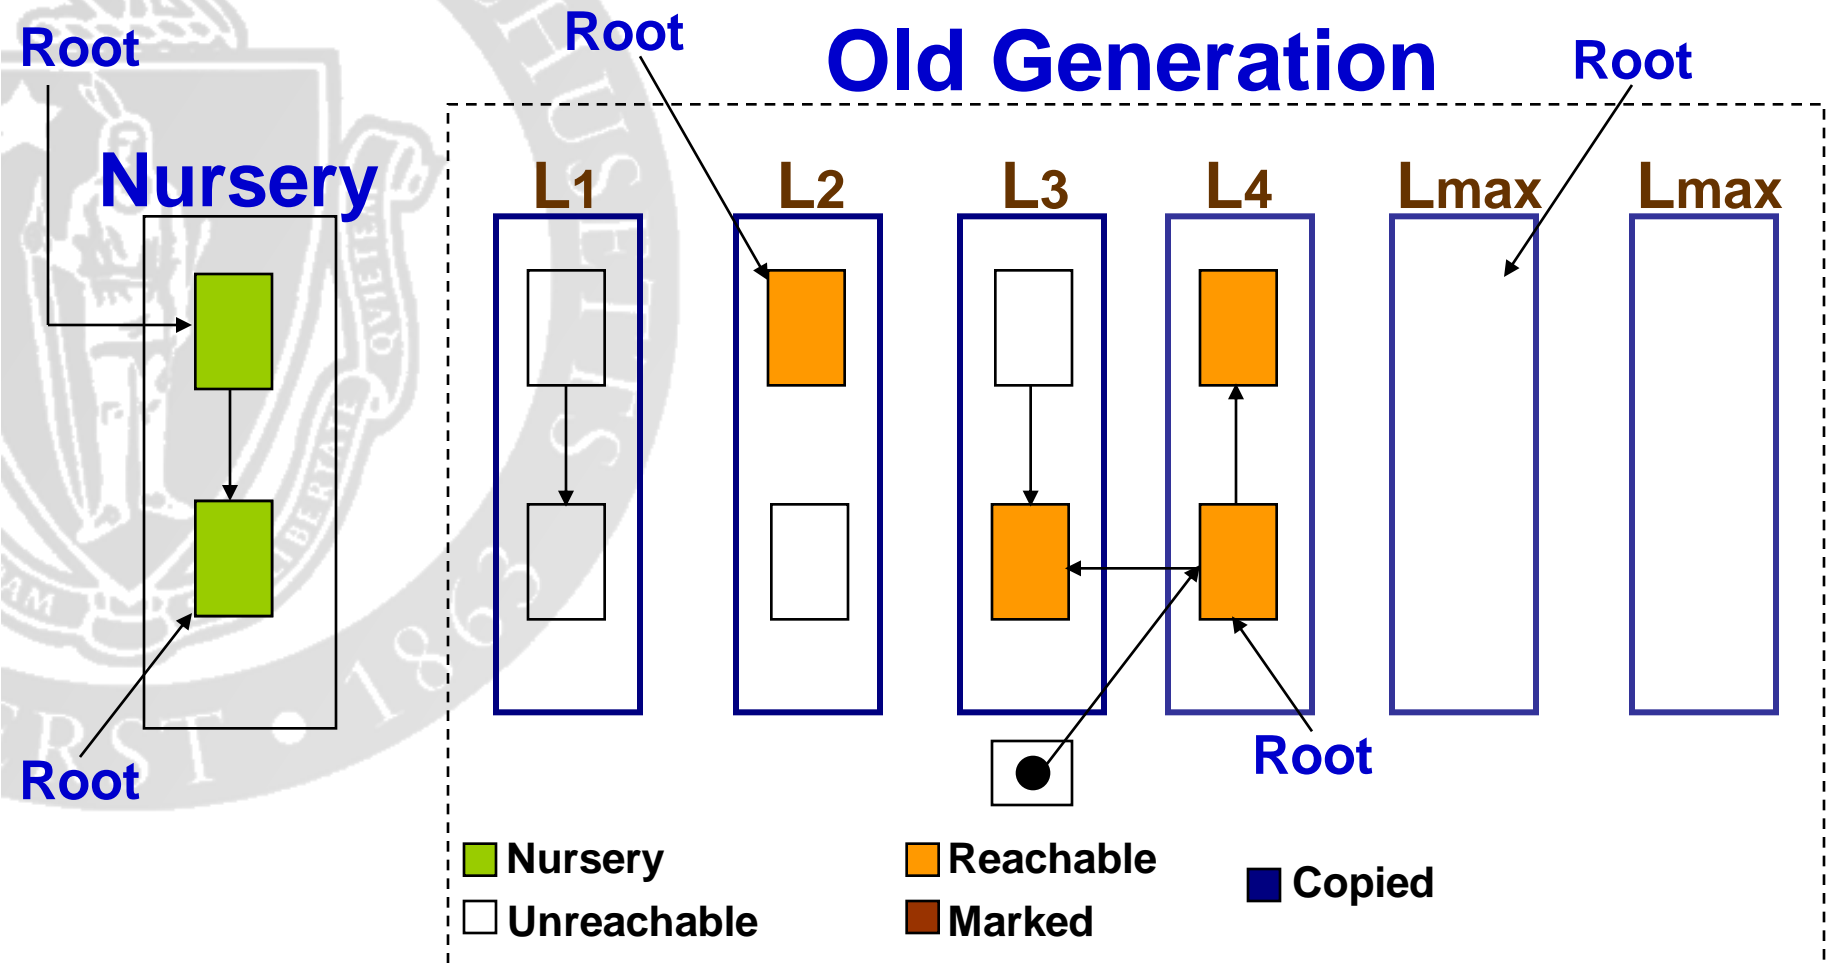


MC² – Mark Phase

- Logically order old gen. windows
- Three mark phase tasks:
 - Mark reachable objects
 - Calculate live data volume in each window
 - Build per-window remembered sets
- Start when old gen. getting full: 80%, say
- Interleave marking with nursery allocation
 - Do some marking after every n bytes allocated
 - Reduces mark phase pauses



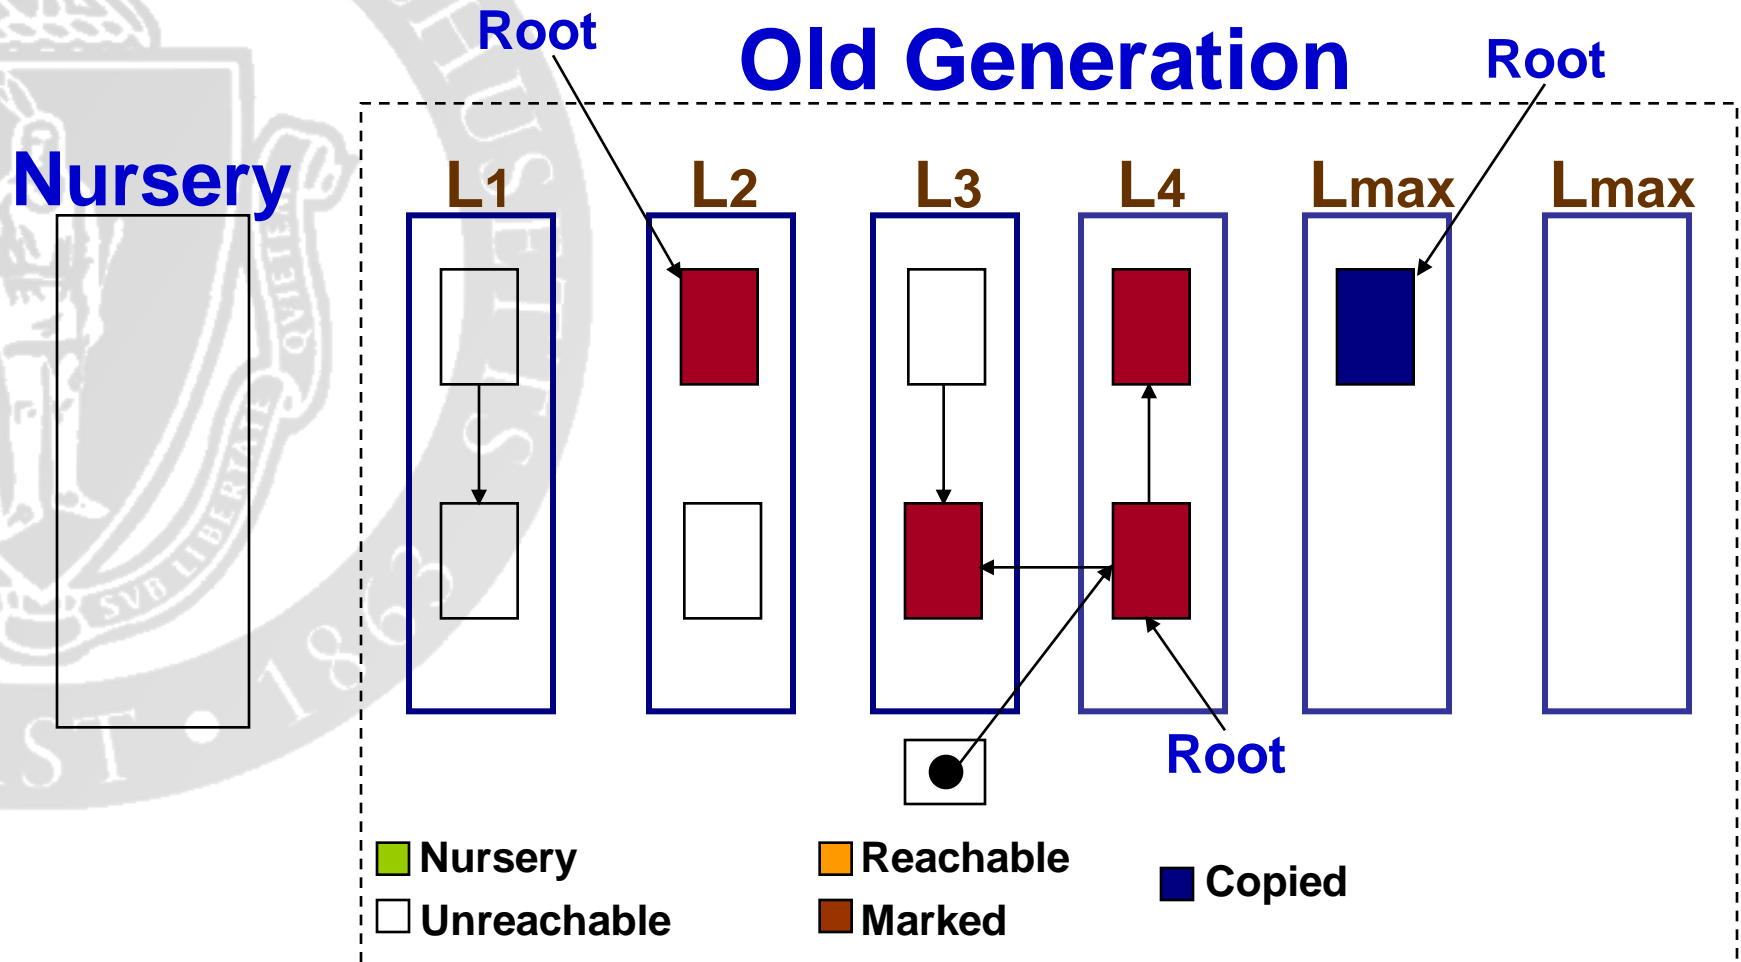
MC² Example – Mark Phase



- MC² marks two objects during nursery allocation



MC² Example – Classify Windows



- Locate high-occupancy windows and discard remsets

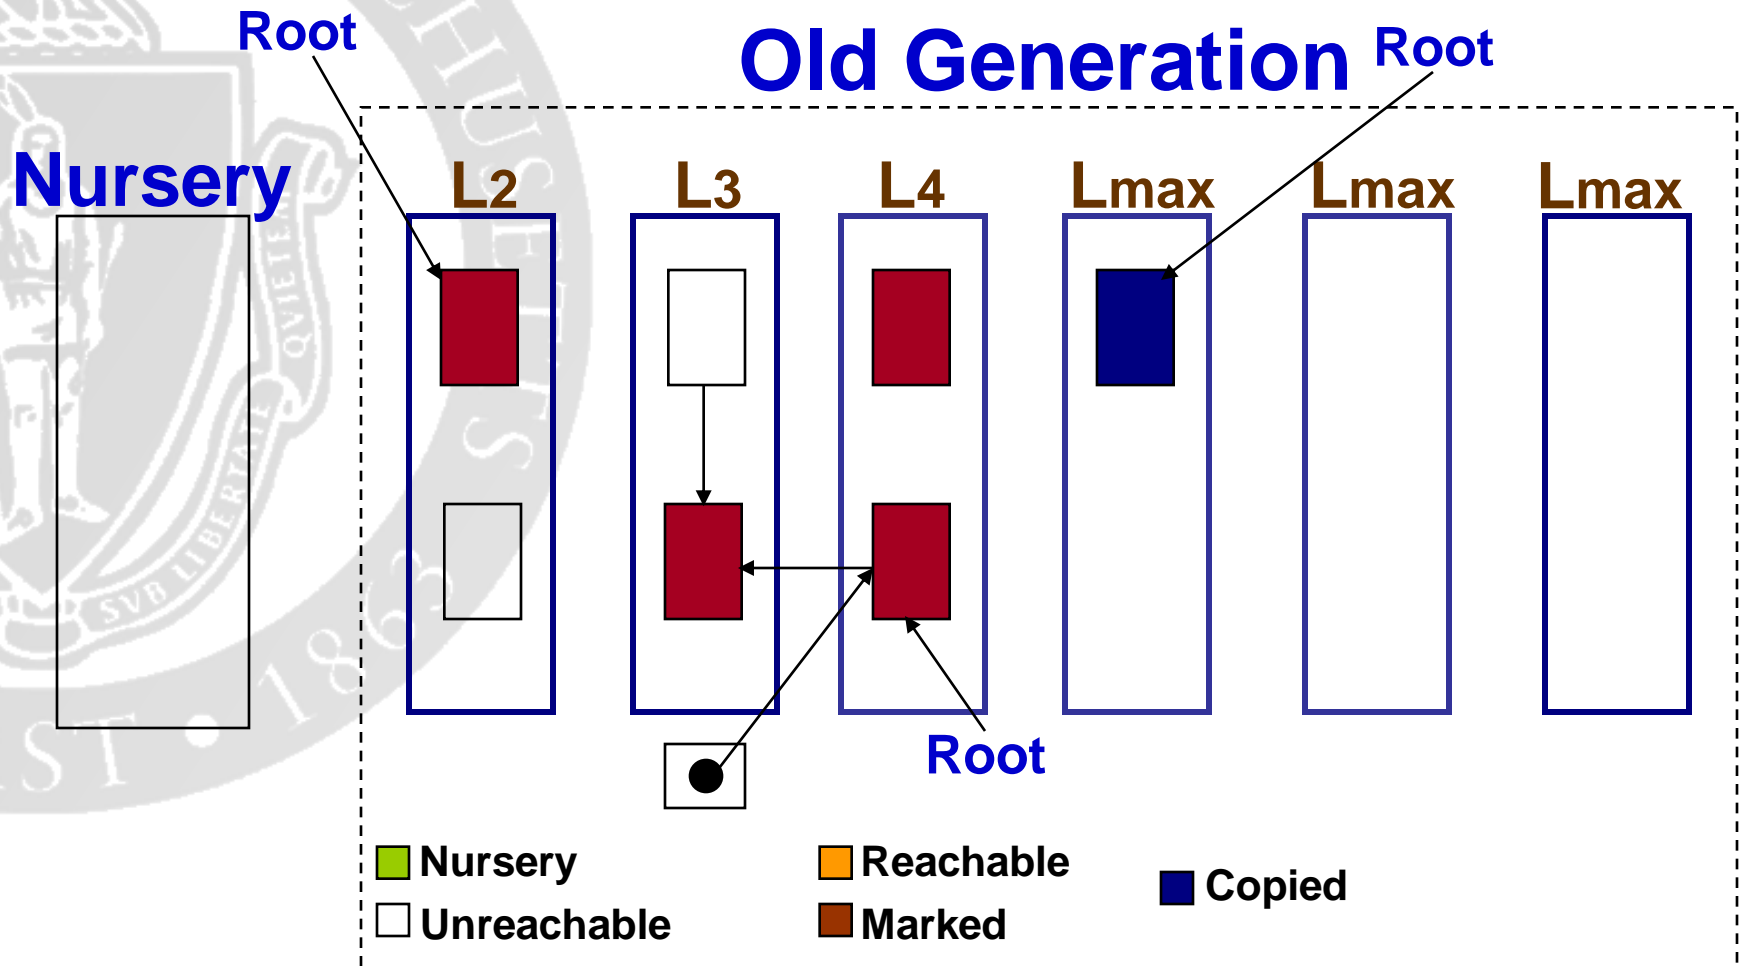


The diagram illustrates a hierarchical tree structure. At the top is the **Root** node. Below it is the **Nursery** level, which contains a large rectangular box. To the right of the Nursery is the **L2** level, which contains a red rectangular box and a white rectangular box. A dashed line separates the Nursery from the L2 level. An arrow points from the Root node to the red box in the L2 level.



- UNIVERSITY OF MASSACHUSETTS AMHERST • Department of Computer Science

MC² Example – Classify Windows



- Locate high-occupancy windows and discard remsets

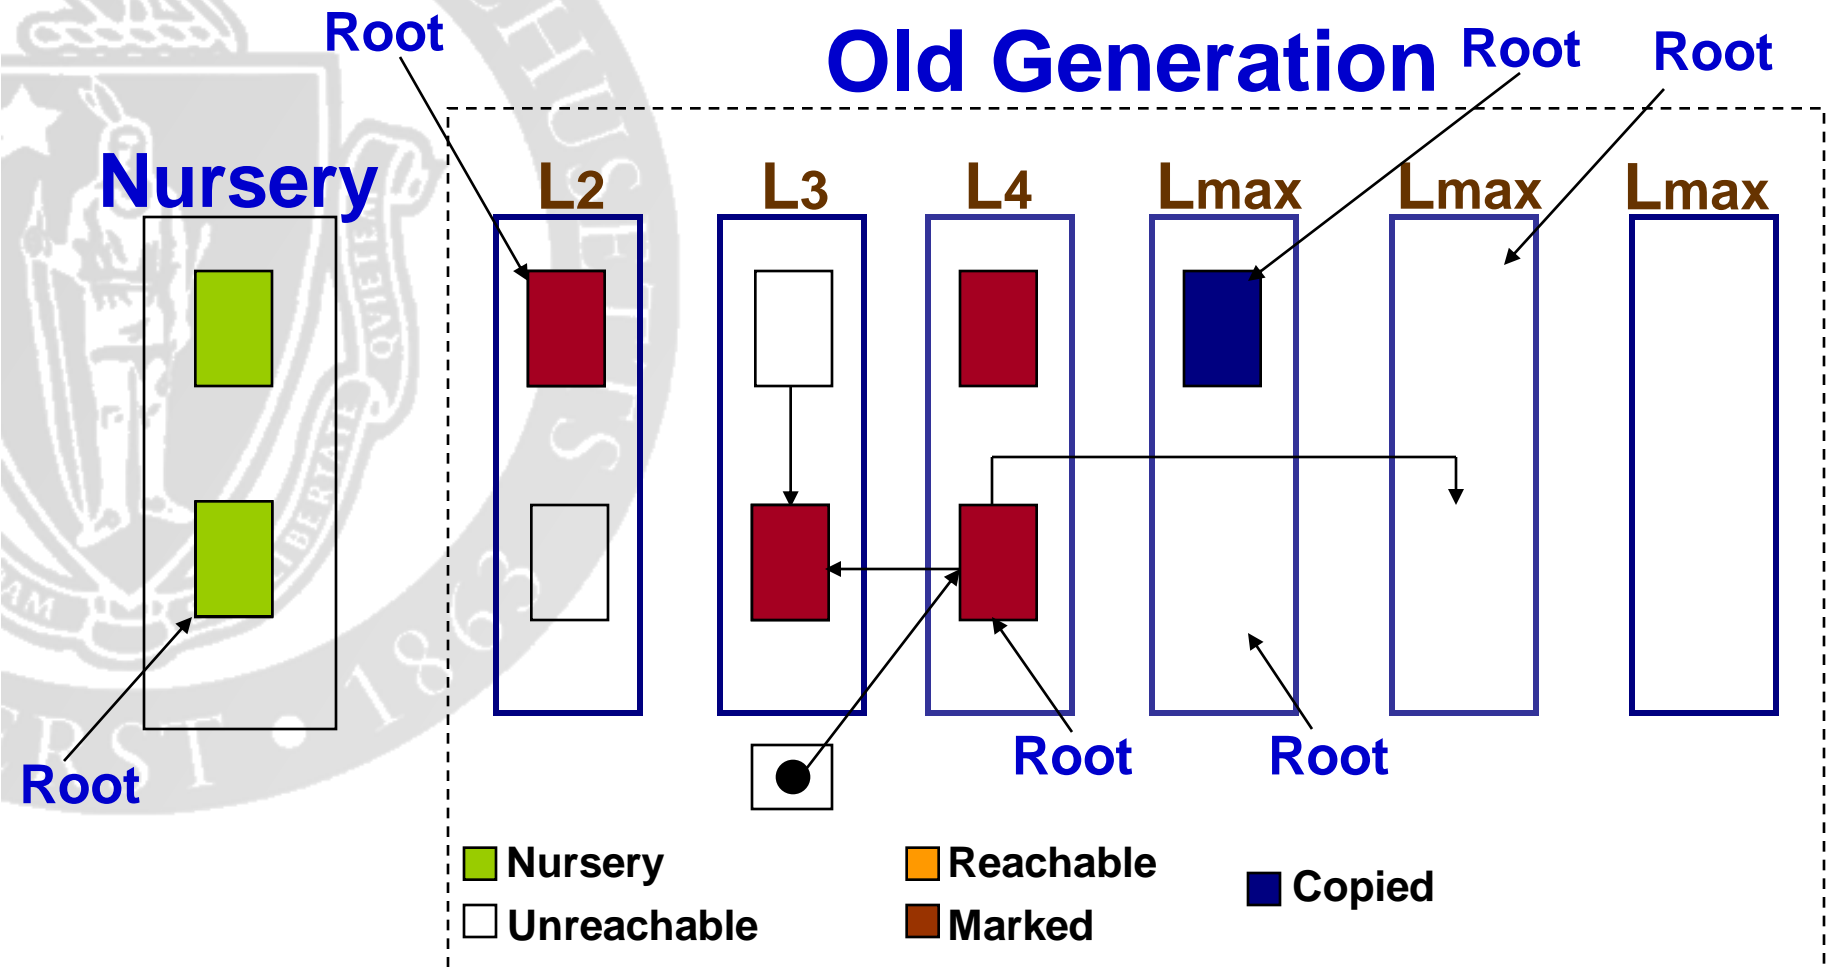


MC² – Copy Phase

- Copy and compact reachable data
- Performed in small increments
 - One windowful of live data copied per increment
- One increment per nursery collection
- High-occupancy windows copied *logically*



MC² Example – Copy Phase



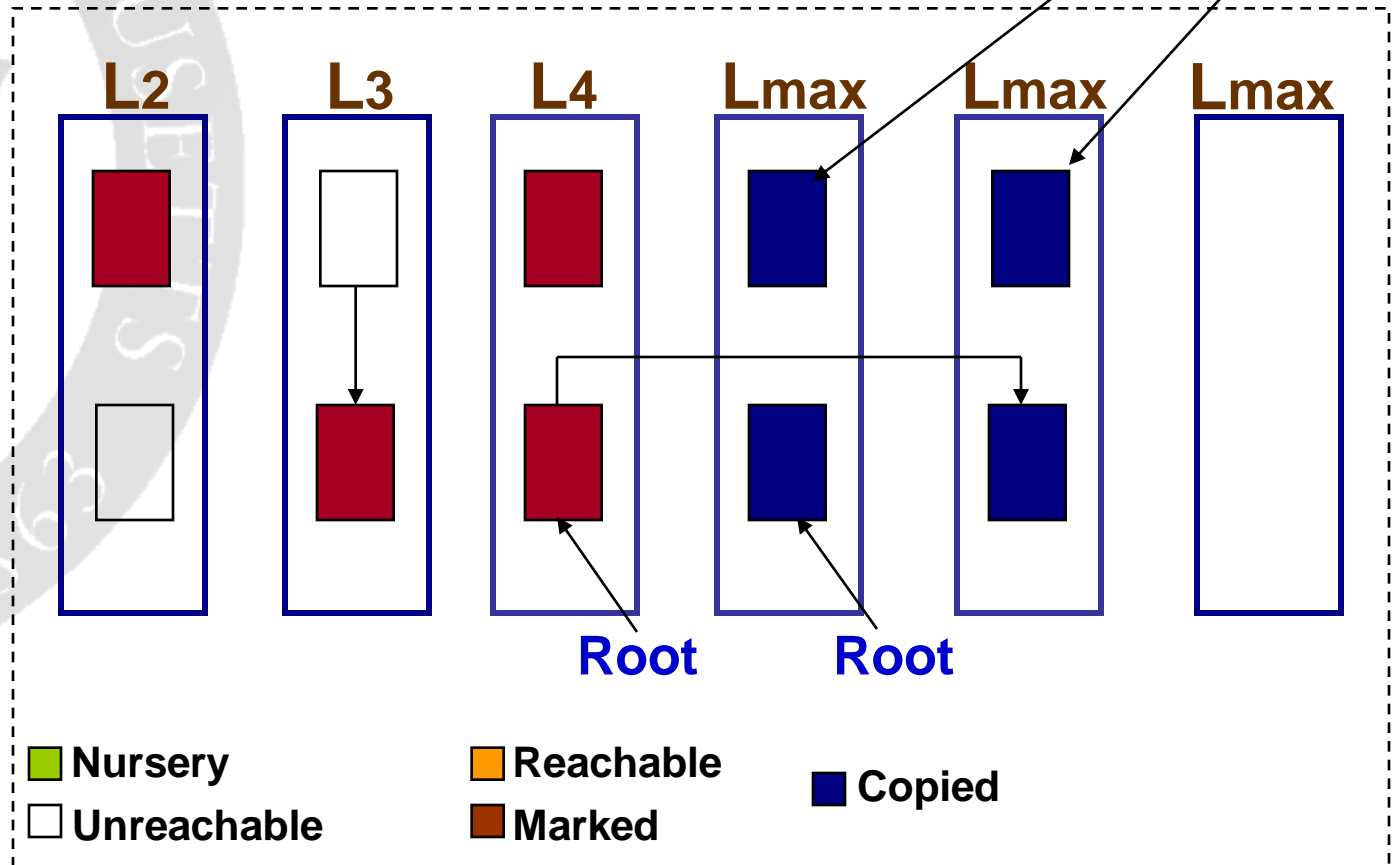
- Copy a window worth of data during nursery collection



MC² Example – Copy Phase

Old Generation

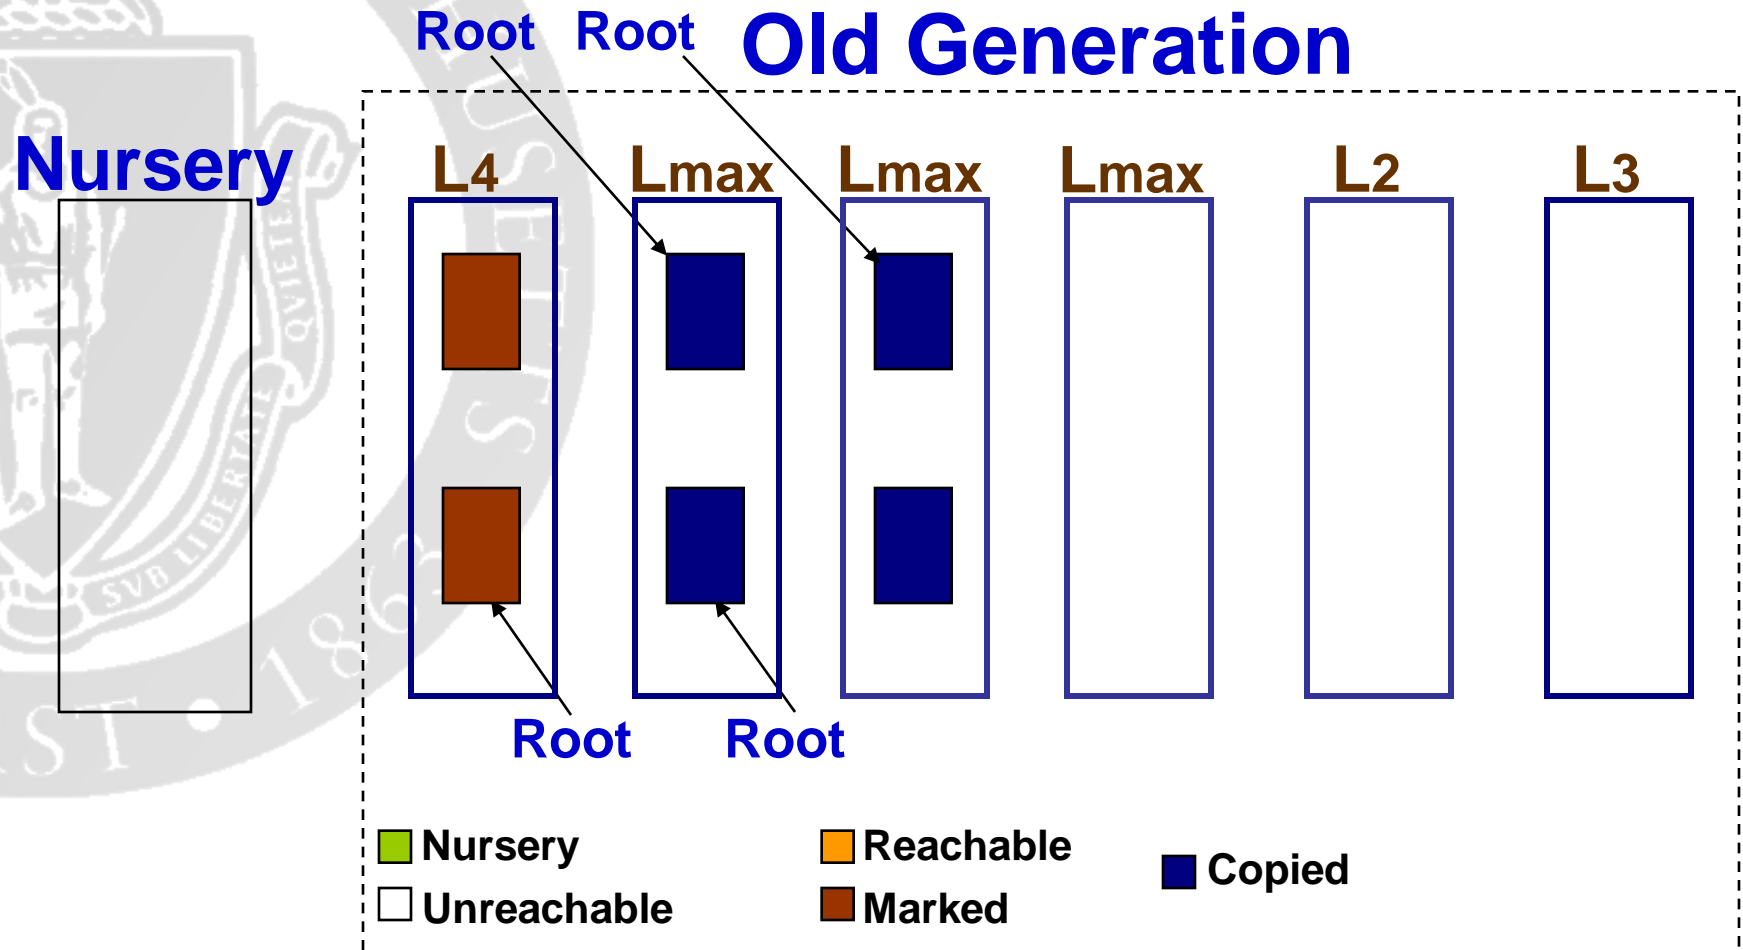
Nursery



- Copy a window worth of data during nursery collection



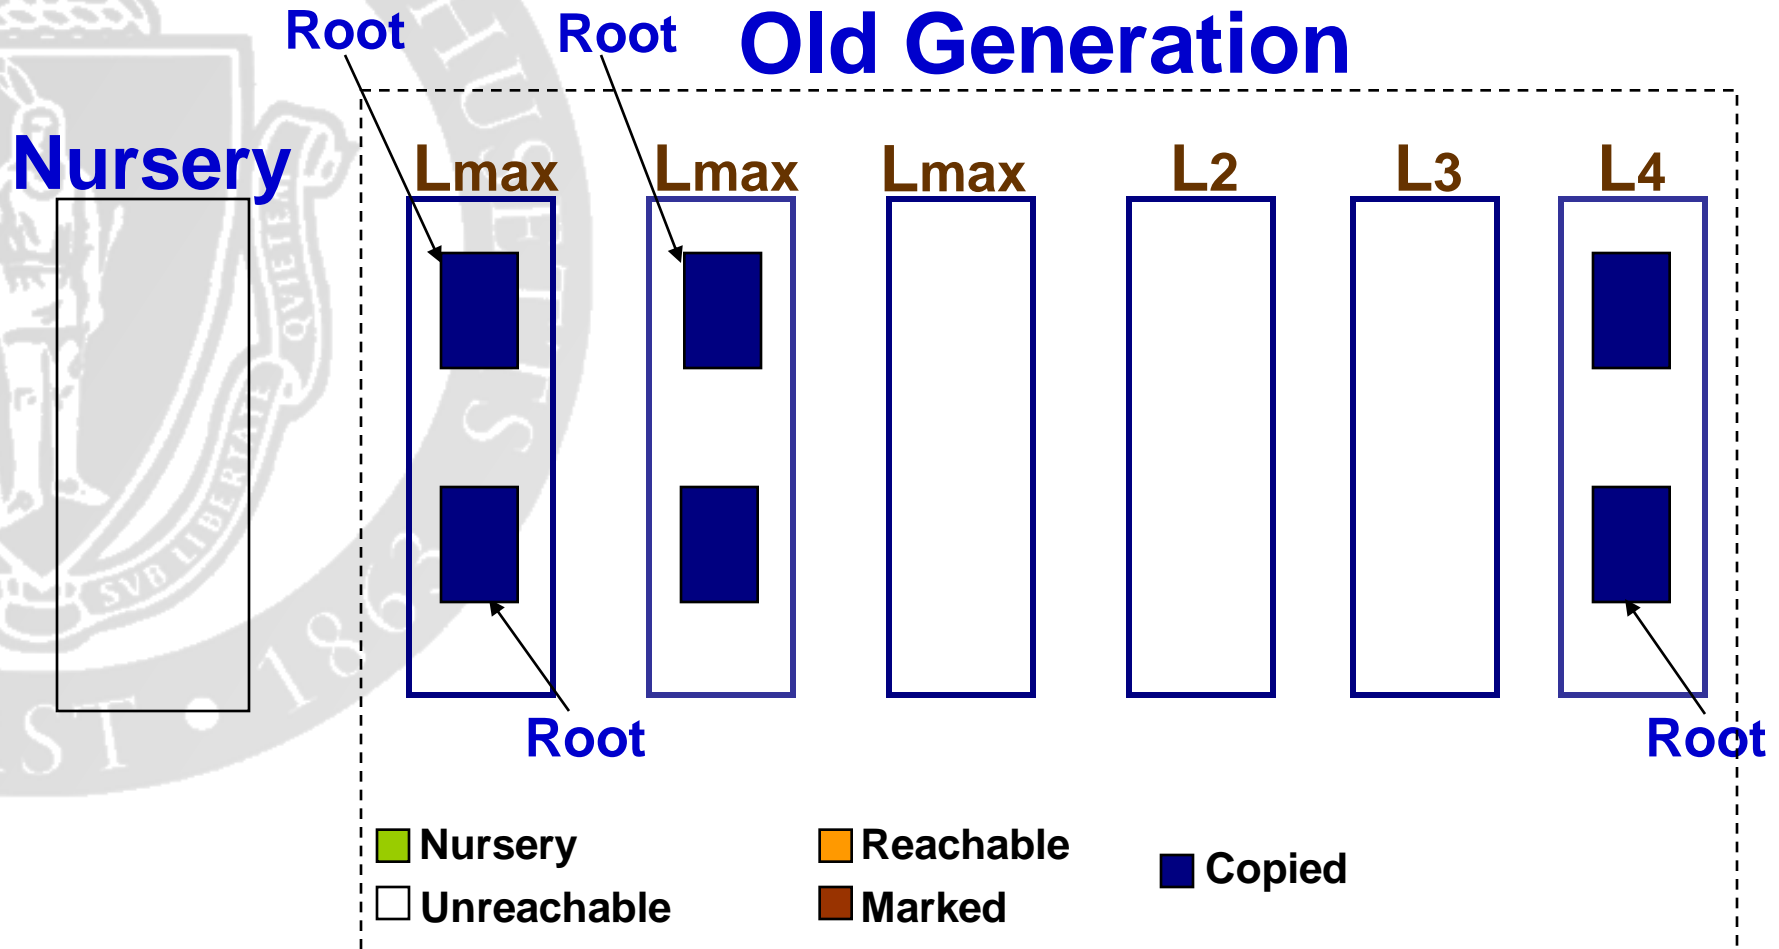
MC² Example – Copy Phase



- Copy a window worth of data during nursery collection



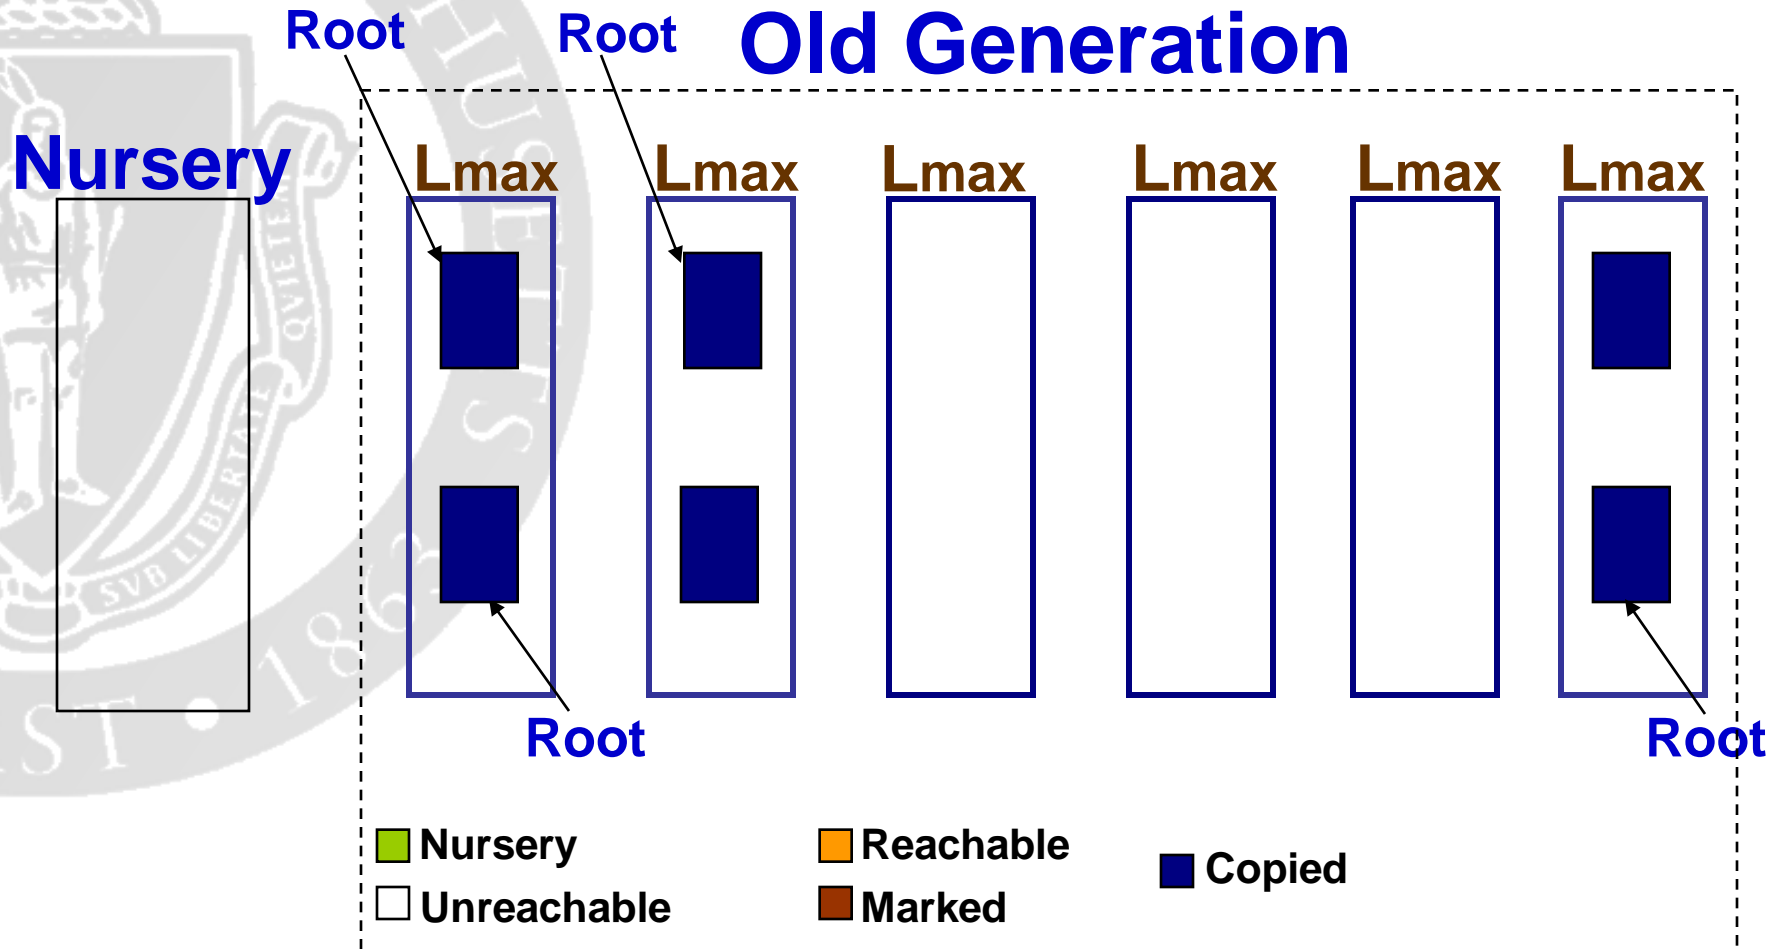
MC² Example – Copy Phase



- Copy a window worth of data during nursery collection



MC² Example – Copy Phase



- Copy a window worth of data during nursery collection




Algorithmic Details

- Large remembered sets
 - Bounding space overhead
- Popular objects
 - Preventing long pauses




Handling large remembered sets



- Normal remembered set for W1 stores 5 pointers (20 bytes on a 32 bit machine)



Handling large remembered sets



- Card table requires only 2 bytes (one per card)




Handling large remembered sets

- Set a limit on the total remembered set size (e.g., 5% of total heap space).
- Replace large remsets with card table when total size approaches limit
- Good tradeoff between speed and space utilization



Handling popular objects


- Divide heap into small *physical* windows
- Normally copy a group of physical windows
- Popular physical windows not grouped



Normal Windows



**Popular
Window**




Normal Windows



Handling popular objects

- Divide heap into small *physical* windows
- Normally copy a group of physical windows
- Popular physical windows not grouped



Handling popular objects

- Identify popular object while converting remset to card table
- Isolate popular object at high end of heap
 - Do not need to maintain references to the objects
- Copying a popular object can cause a long pause, but does not recur

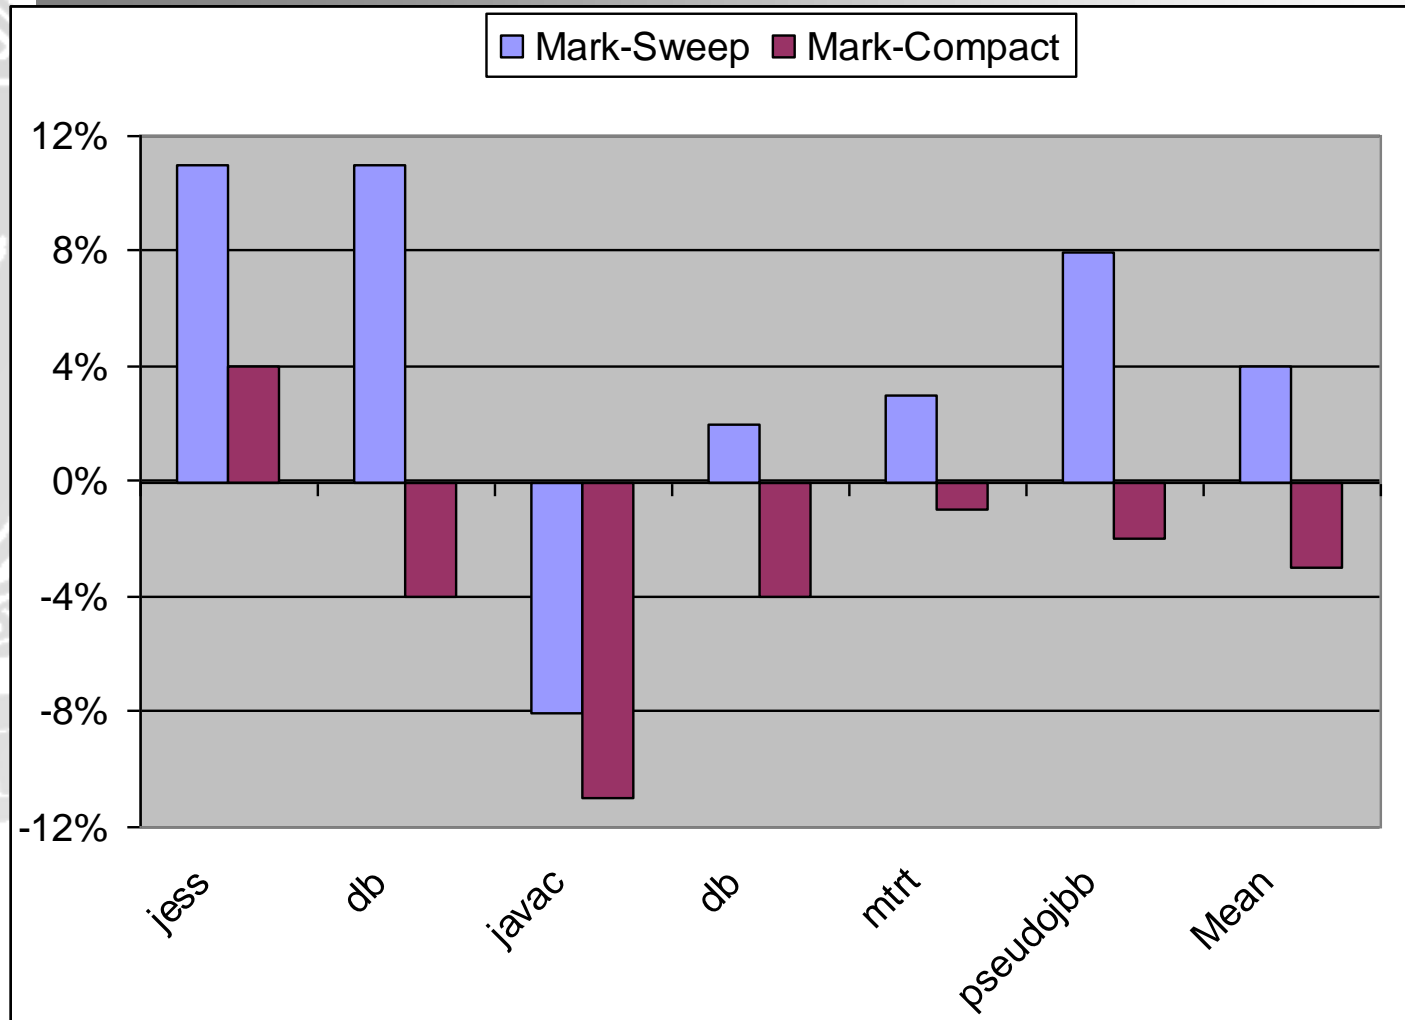


Experimental Results

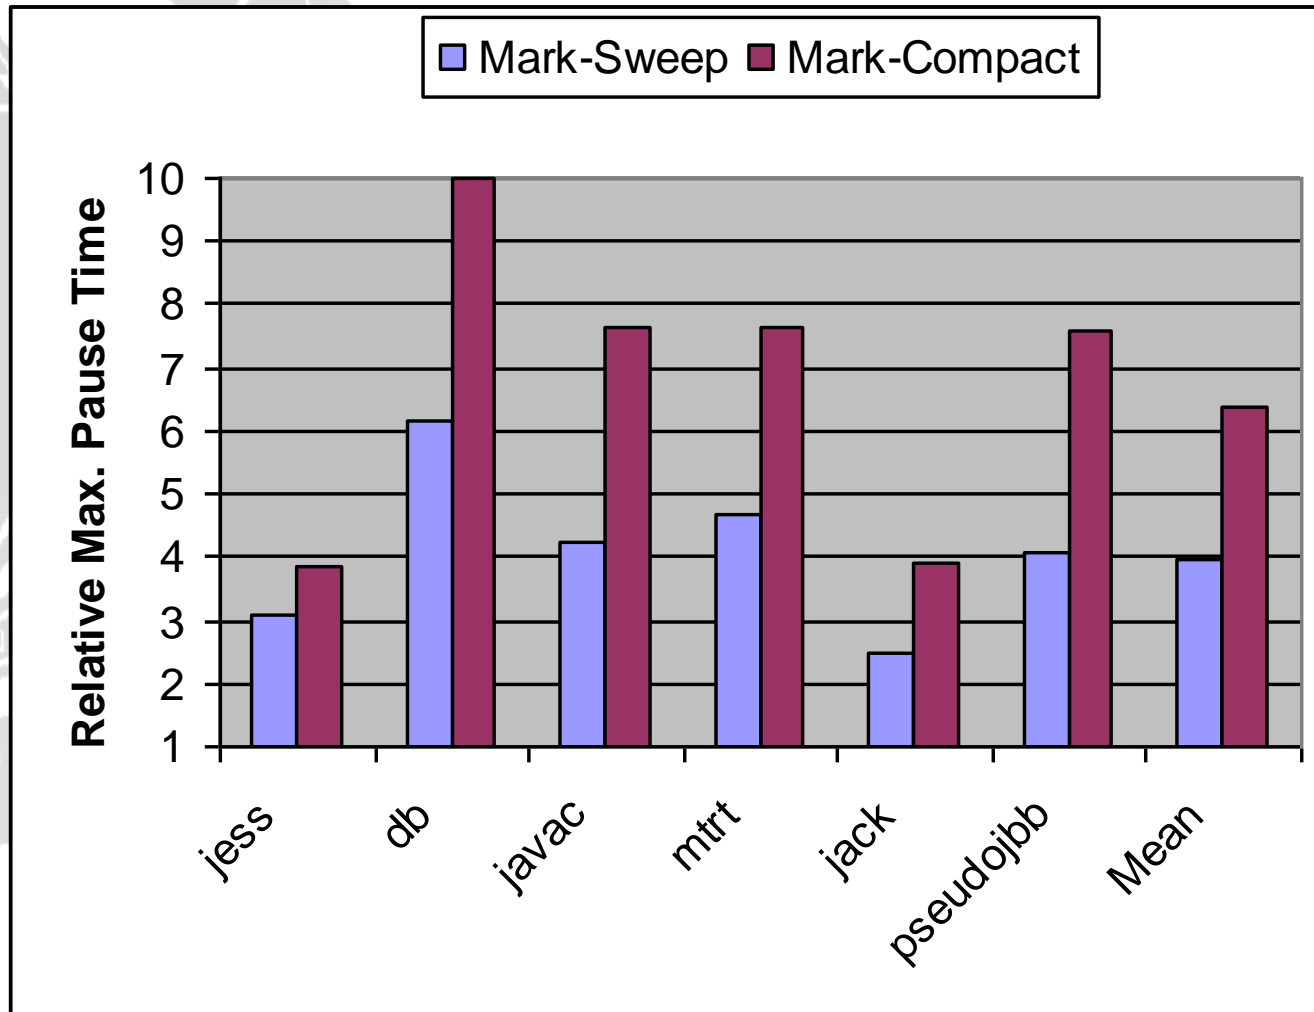
- Implemented in Jikes RVM (2.2.3)/MMTk
- Pentium 4 1.7 GHz, 512MB memory, RedHat Linux 2.4.7-10
- Benchmarks: SPECjvm98, pseudojbb
- Collectors evaluated
 - Generational Mark-Sweep (MS)
 - Generational Mark-(Sweep)-Compact (MSC)
 - MC²
- MSC, MC² use separate code and data spaces
- Results: Execution time, pause time in a heap
1.8x program live size



Execution Time relative to MC² (Heap Size = 1.8x max. live size)



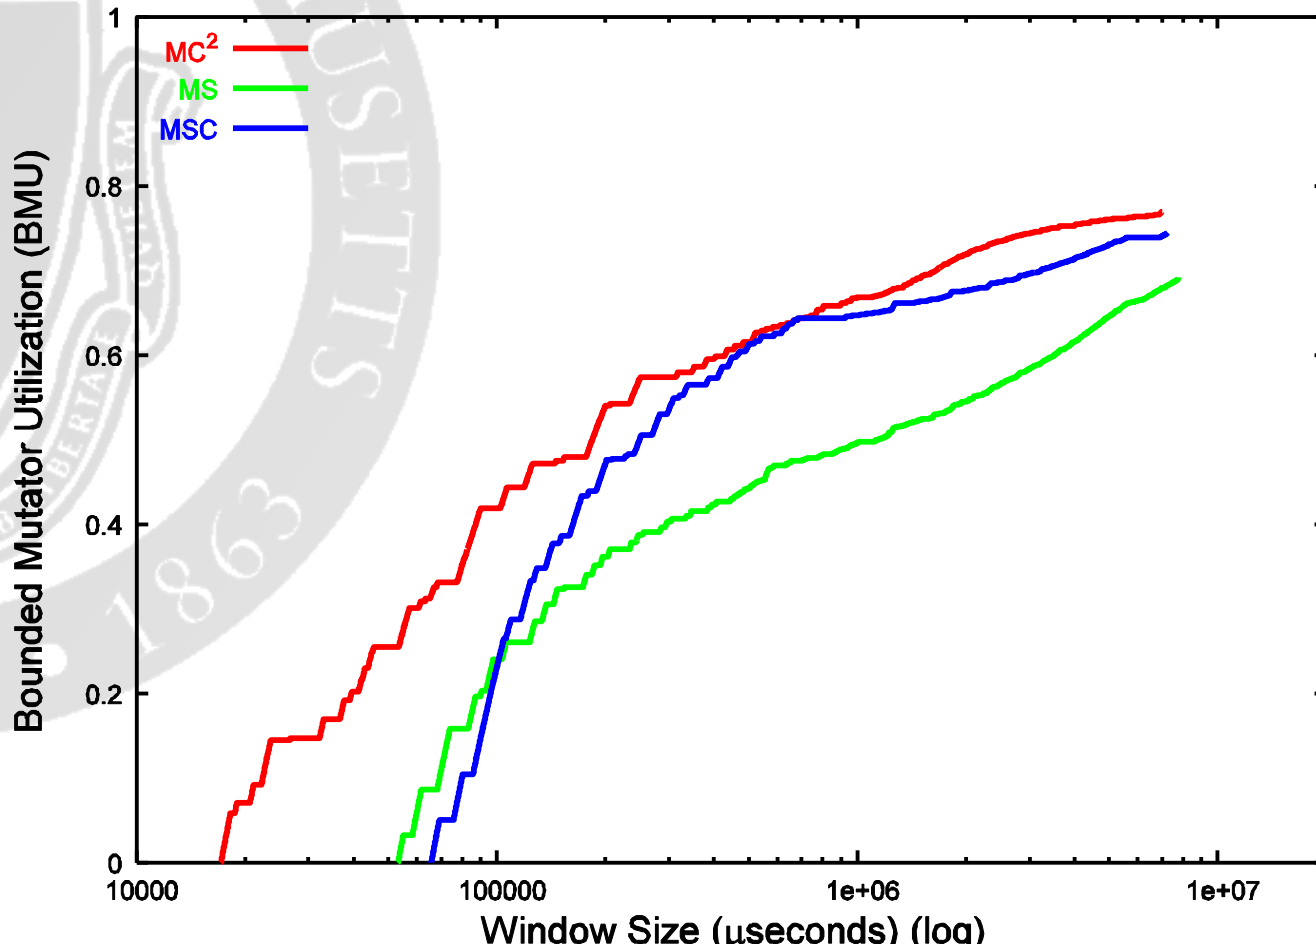
Max. Pause Time relative to MC² (Heap Size = 1.8x max. live size)




- MC² max. pause range: 17-41ms



Bounded Mutator Utilization (jess)



Bounded Mutator Utilization (javac)



Conclusions

- MC²: suitable for handheld devices with soft real time requirements
 - Low space overhead (50-80%)
 - Good throughput (3-4% slower than non-incremental compacting collector)
 - Short pause times (17-41ms, factor of 6 lower than non-incremental compacting collector)
 - Well distributed pauses
- Also suitable for desktop environments



Backup Slides



Traditional Tracing Collectors

- Mark-Sweep
 - Fragmentation
 - Locality effects
- Mark-Compact
 - Long pauses
- Copying
 - 2X space overhead

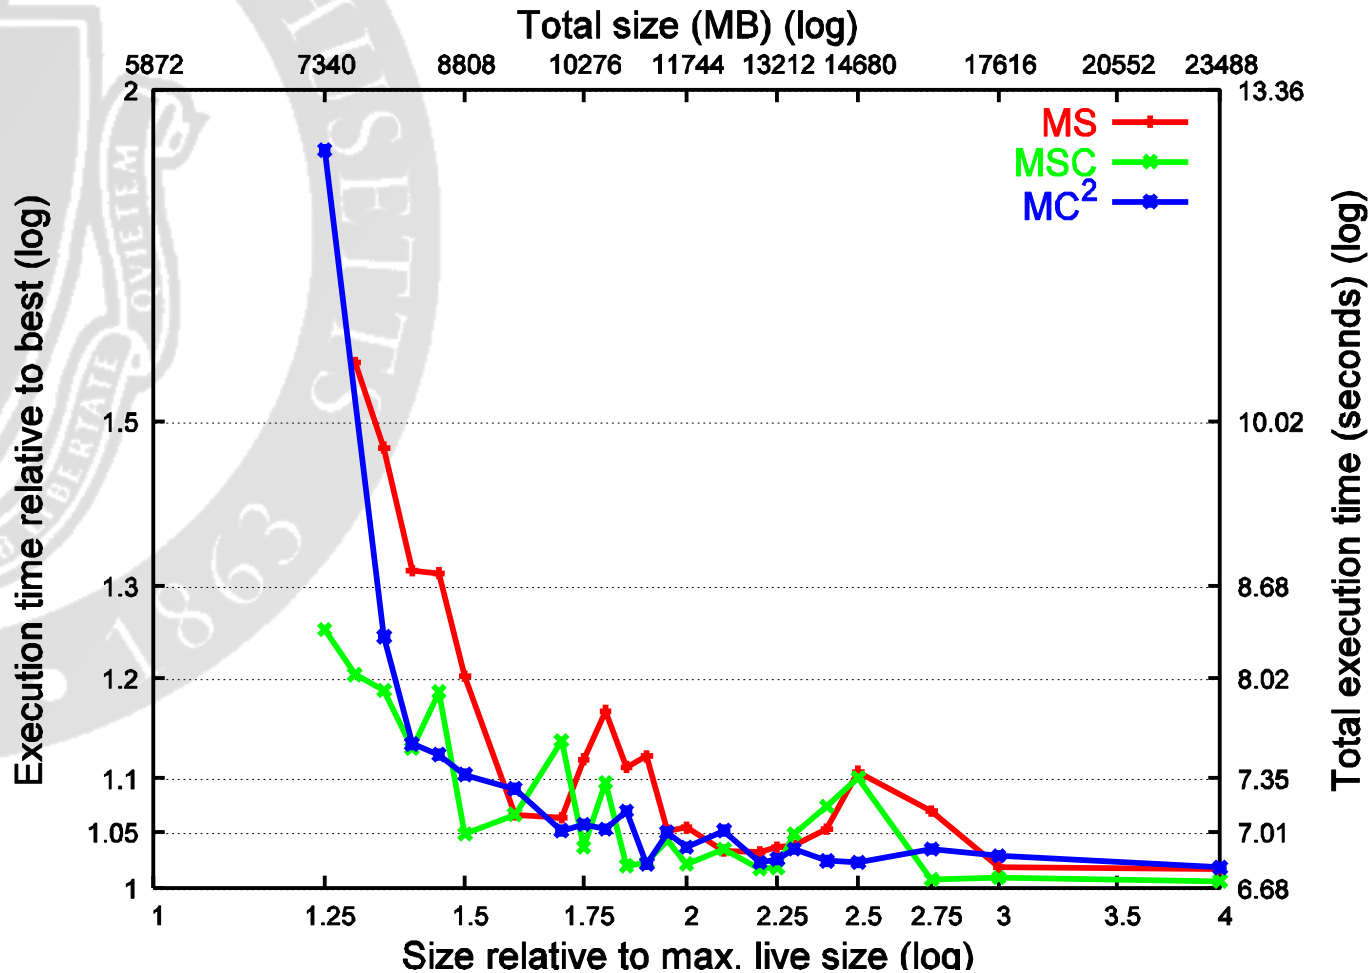


Modern incremental collectors

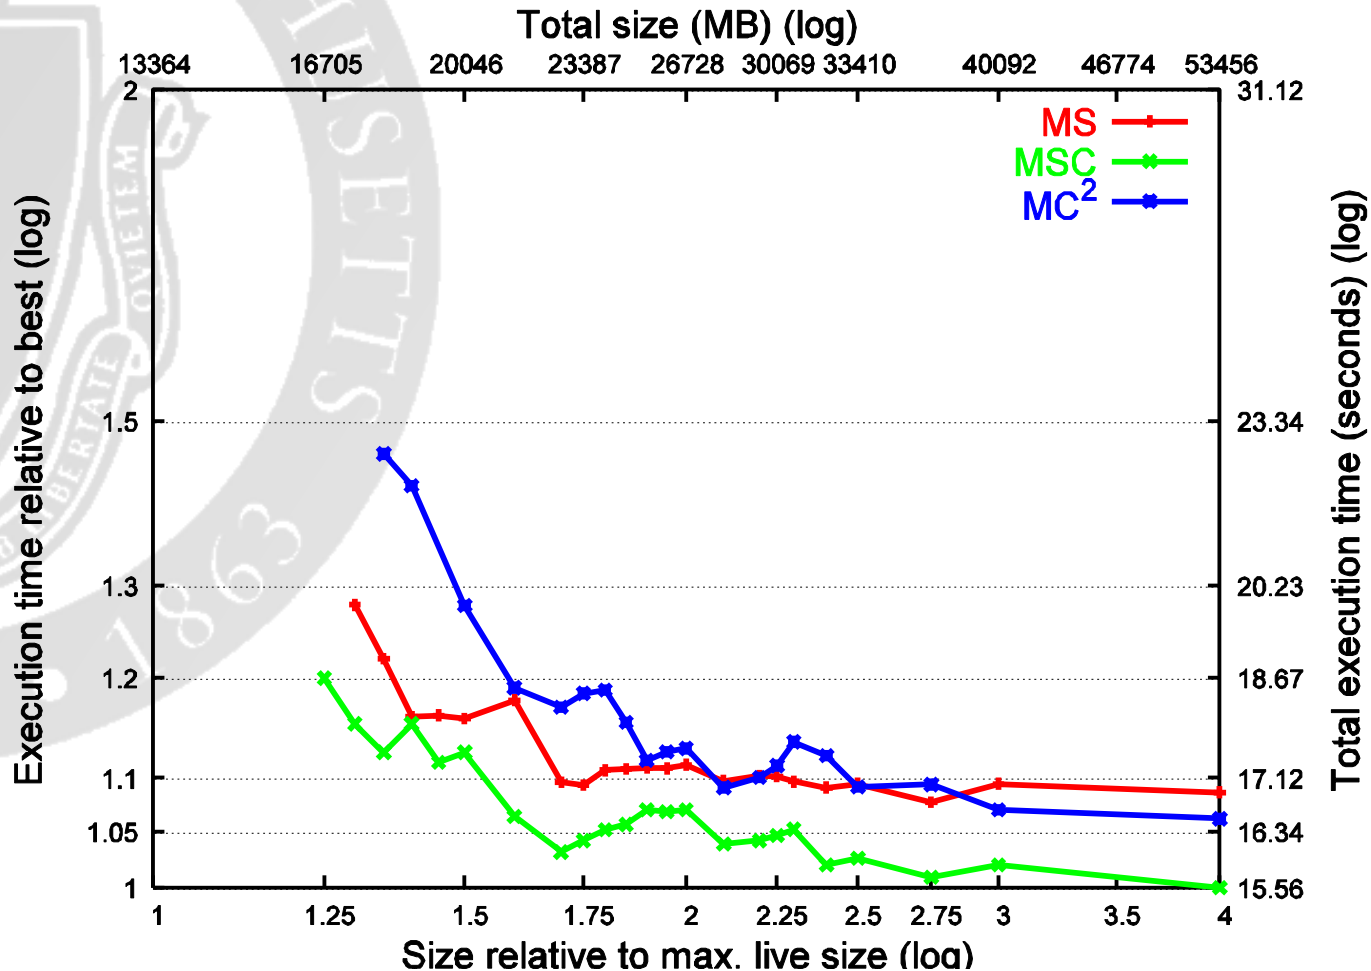
- Train collector – Hudson and Moss
 - Performs well in moderate size heaps
 - High copying overhead
- Lang and Dupont, Ben Yitzhak et al
 - Good throughput, pause times
 - Do not address metadata overheads
- Bacon, Cheng, Rajan
 - Address memory-constrained device requirements
 - Require advanced compiler optimizations



jess Execution Time



javac Execution Time

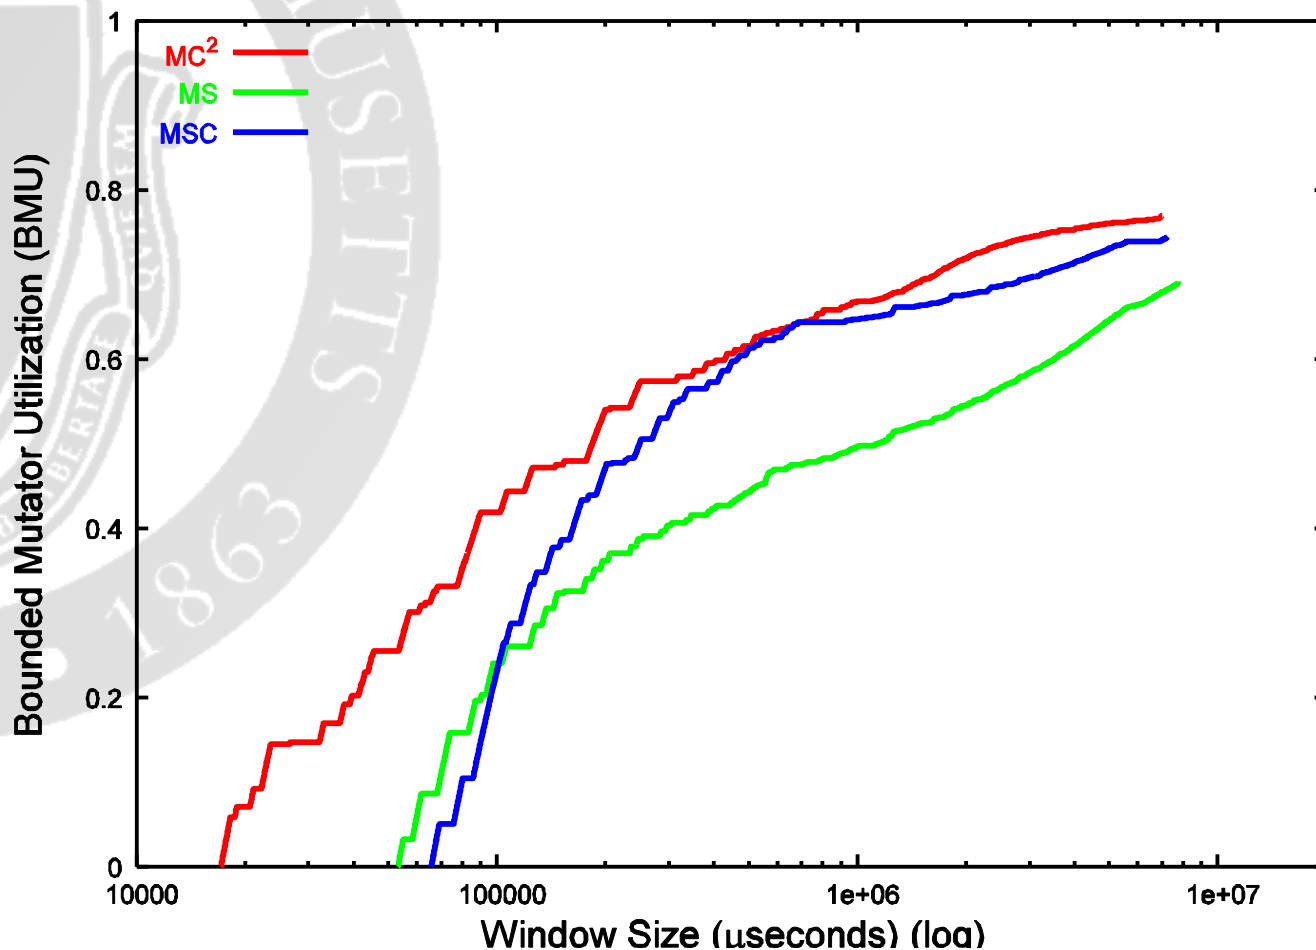


Pause Time, Execution Time Comparison (80% space overhead)

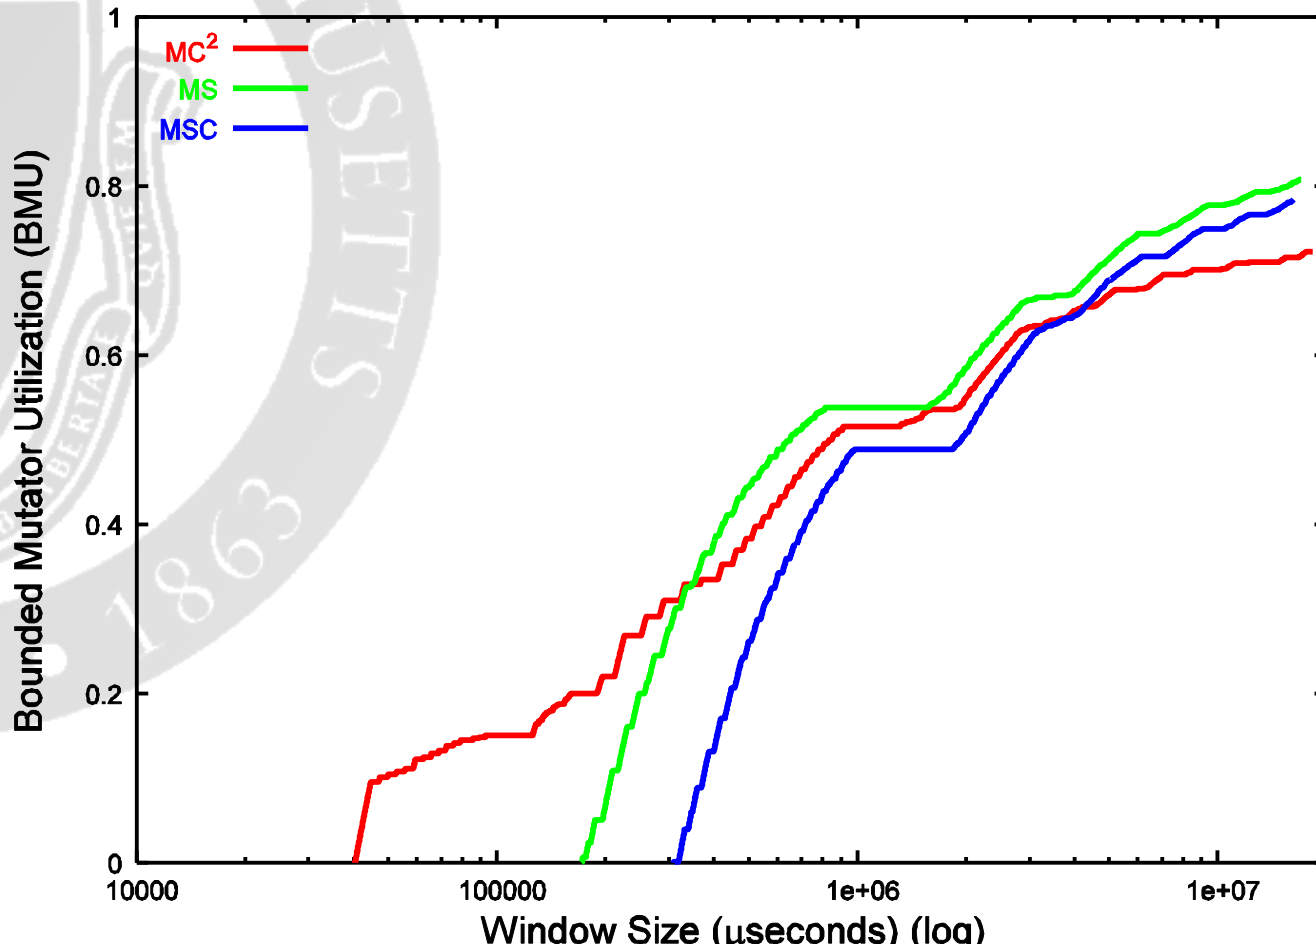
Benchmark	MC ² MPT (ms)	MS MPT (ms)	MS/MC ² ET	MSC MPT (ms)	MSC/MC ² ET
_202_jess	17.2	53.2	1.11	65.7	1.04
_209_db	19.9	123.0	1.11	198.5	0.96
_213_javac	40.4	171.9	0.92	308.9	0.89
_227_mtrt	29.6	138.1	1.02	225.2	0.96
_228_jack	23.9	59.7	1.03	92.9	0.99
pseudojbb	41.5	168.2	1.08	314.5	0.98
Geo. Mean	27.2	107.7	1.04	172.7	0.97



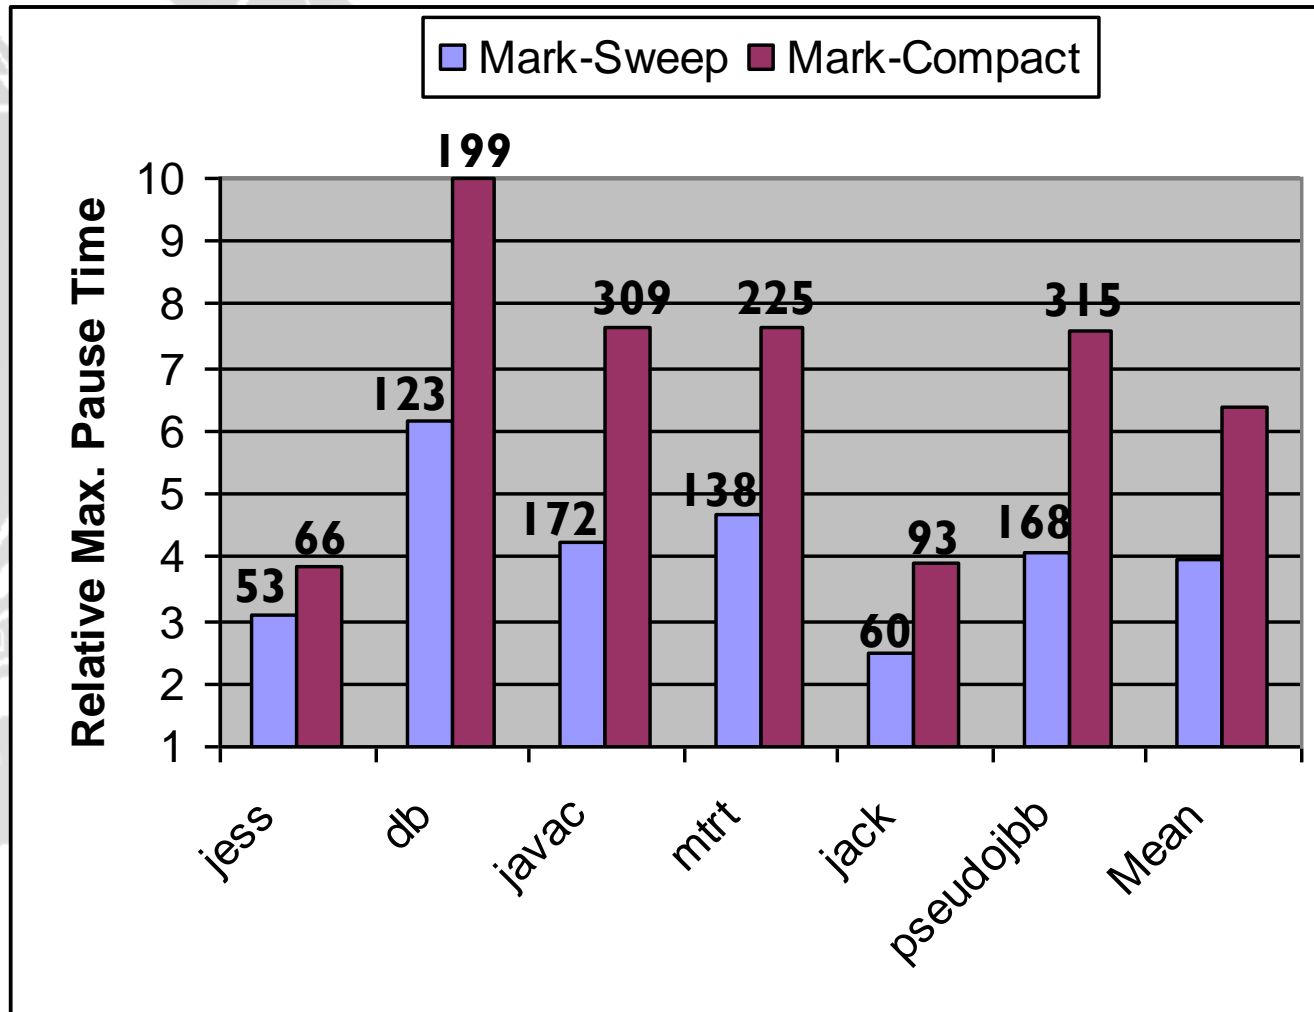
Bounded Mutator Utilization (jess)



Bounded Mutator Utilization (javac)




Max. Pause Time relative to MC² (Heap Size = 1.8x max. live size)



- MC² max. pause range: 17-41ms



Incremental Marking Error



Handling marking error

- Track mutations using write barrier
- Record modified old generation objects
- Scan these modified objects at GC time
- Record interesting slots in remembered sets

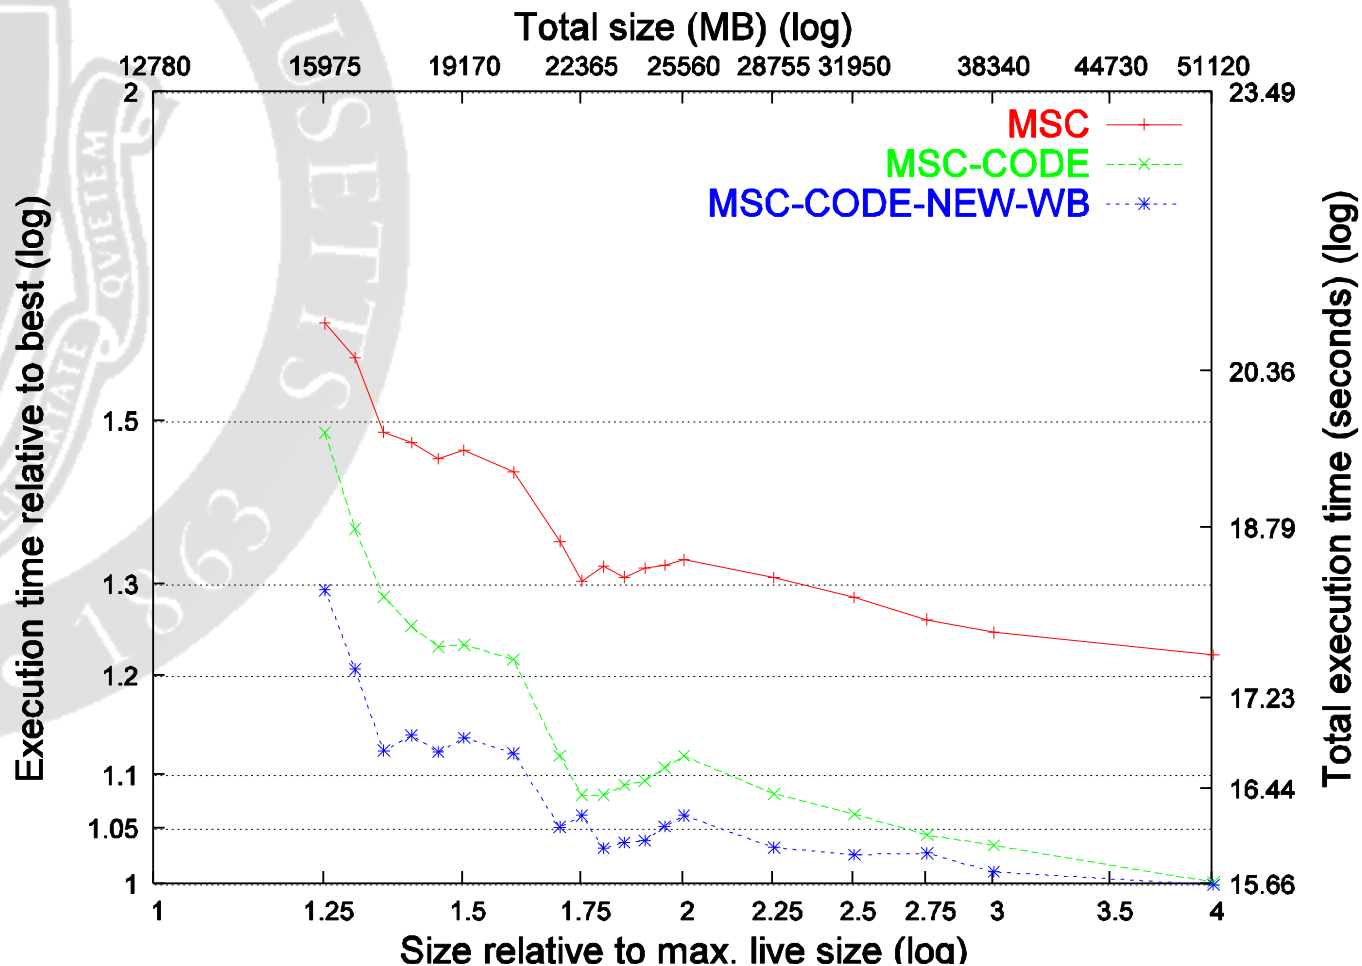


JMTk Heap Layout

Boot Image	Imm. Object Space	Large Object Space	Old Generation	Nursery
---------------	-------------------------	--------------------------	----------------	---------




_213_javac (MSC execution time)



MC² Overview

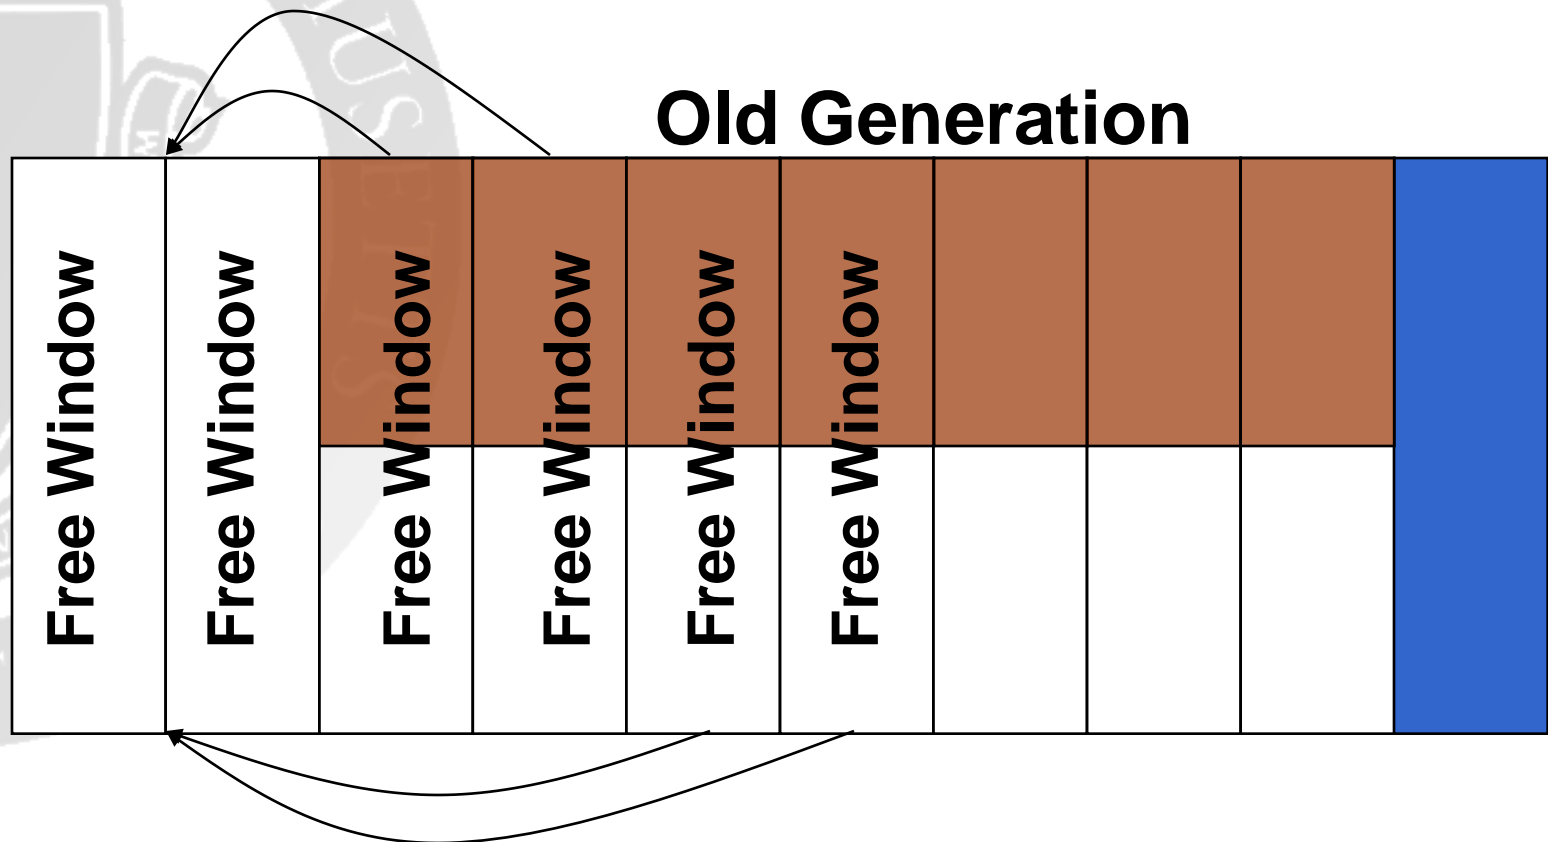
Old Generation



■ Reachable data □ Unreachable data ■ Copied data



MC² Overview




■ Reachable data □ Unreachable data ■ Copied data



MC² Overview

Old Generation



■ Reachable data □ Unreachable data ■ Copied data

