



Reconsidering Custom Memory Allocation

Emery Berger, Ben Zorn, Kathryn McKinley




Microsoft
Research

THE UNIVERSITY OF
T E X A S
— AT AUSTIN —



Custom Memory Allocation

- Programmers replace new/delete, bypassing system allocator
 - Reduce runtime – *often*
 - Expand functionality – *sometimes*
 - Reduce space – *rarely*
- **Very common practice**
 - Apache, gcc, lcc, STL, database servers...
 - Language-level support in C++
 - Widely recommended



“Use custom allocators”



Drawbacks of Custom Allocators

- Avoiding system allocator:
 - More code to maintain & debug
 - Can't use memory debuggers
 - Not modular or robust:
 - Mix memory from custom and general-purpose allocators → **crash!**
- ⇒ **Increased burden on programmers**

Are custom allocators really a win?




Overview

- Introduction
- Perceived benefits and drawbacks
- Three main kinds of custom allocators
- Comparison with general-purpose allocators
- Advantages and drawbacks of regions
- Reaps – generalization of regions & heaps

(I) Per-Class Allocators

- Recycle freed objects from a free list


```
a = new Class1;  
b = new Class1;  
c = new Class1;  
delete a;  
delete b;  
delete c;  
a = new Class1;  
b = new Class1;  
c = new Class1;
```



- + *Fast*
 - + Linked list operations
- + *Simple*
 - + Identical semantics
 - + C++ language support
- *Possibly space-inefficient*

(II) Custom Patterns

- Tailor-made to fit allocation patterns
 - Example: *197.parser* (natural language parser)



```
a = xalloc(8);  
b = xalloc(16);  
c = xalloc(8);  
xfree(b);  
xfree(c);  
d = xalloc(8);
```

+ *Fast*

+ Pointer-bumping allocation

■ *Brittle*

■ Fixed memory size

■ Requires stack-like lifetimes

(III) Regions

- Separate areas, deletion only *en masse*

regioncreate(r)
regionmalloc(r, sz)
regiondelete(r)



- + *Fast*
 - + Pointer-bumping allocation
 - + Deletion of chunks
- + *Convenient*
 - + One call frees all memory
- *Risky*
 - Dangling references
 - Too much space


- Increasingly popular custom allocator



Overview


- Introduction
- Perceived benefits and drawbacks
- Three main kinds of custom allocators
- Comparison with general-purpose allocators
- Advantages and drawbacks of regions
- Reaps – generalization of regions & heaps

Custom Allocators Are Faster...



- As good as and sometimes *much faster* than Win32

Not So Fast...



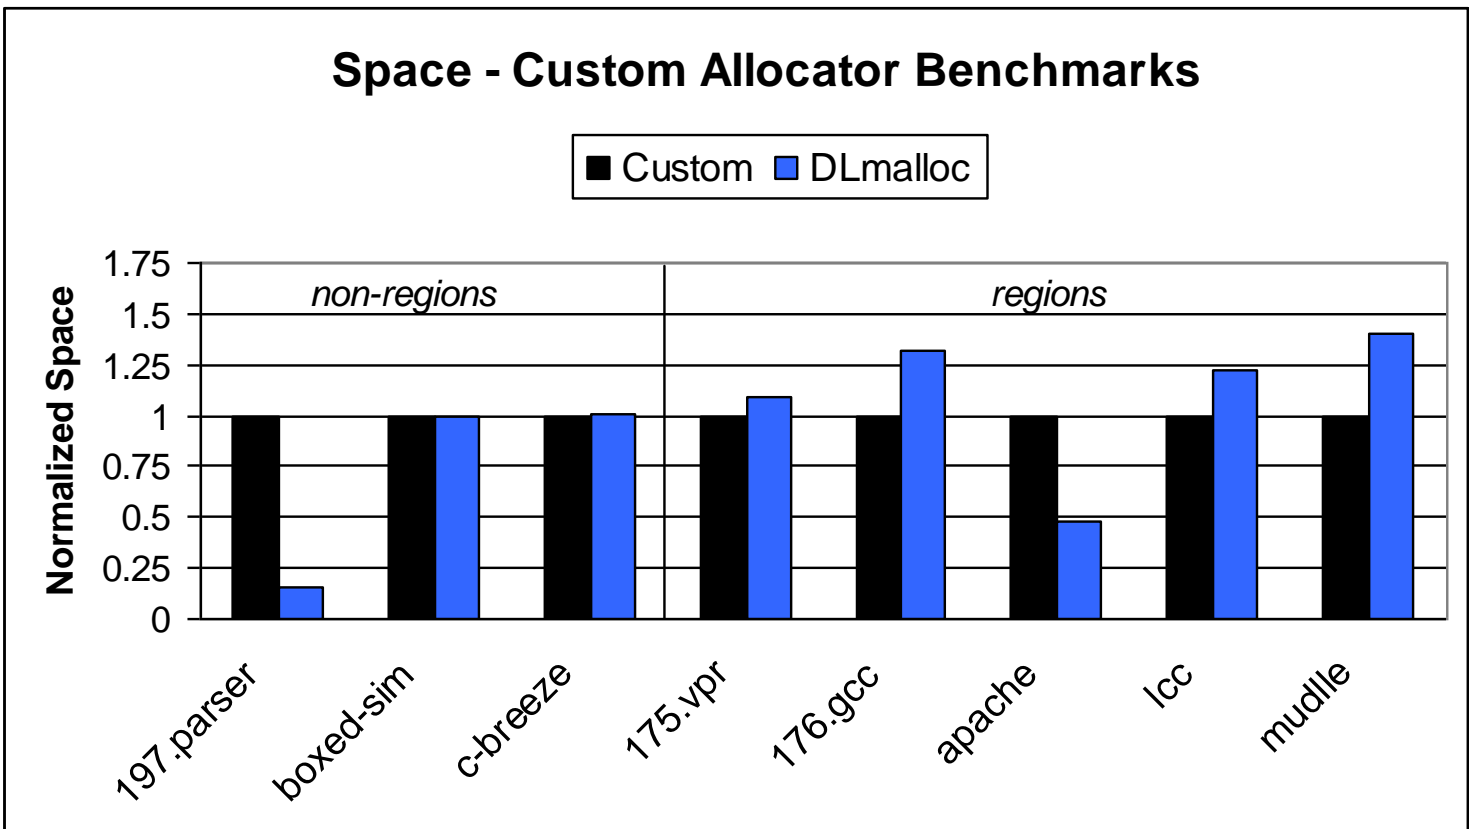
- DLmalloc: as fast or *faster* for most benchmarks



The Lea Allocator (DLmalloc 2.7.0)

- Mature public-domain general-purpose allocator
- Optimized for common allocation patterns
 - Per-size quicklists \approx per-class allocation
- Deferred coalescing
(combining adjacent free objects)
 - ⇒ Highly-optimized fastpath
- Space-efficient

Space Consumption: Mixed Results





Overview

- Introduction
- Perceived benefits and drawbacks
- Three main kinds of custom allocators
- Comparison with general-purpose allocators
- Advantages and drawbacks of regions
- Reaps – generalization of regions & heaps



Regions – Pros and Cons

- + Fast, convenient, etc.
- + Avoid resource leaks (e.g., Apache)
 - Tear down memory for terminated connections
- No individual object deletion
 - ⇒ **Unbounded memory consumption**
(producer-consumer, long-running computations, off-the-shelf programs)
 - Apache: vulnerable to DoS, memory leaks

Reap Hybrid Allocator


- $\text{Reap} = \text{region} + \text{heap}$
 - Adds individual object deletion & heap

reapcreate(r)

reapmalloc(r, sz)


reapfree(r,p)

reapdelete(r)




- + Can reduce memory consumption
- + Fast
 - + Adapts to use (region or heap style)
 - + Cheap deletion

Reap Runtime



Reap Space

Space - Custom Allocator Benchmarks





Reap: Best of Both Worlds

- Allows mixing of regions and new/delete
- Case study:
 - New Apache module “*mod_bc*”
 - *bc*: C-based arbitrary-precision calculator
 - Changed 20 lines out of 8000
 - Benchmark: compute 1000th prime
 - With Reap: 240K
 - Without Reap: 7.4MB



Conclusions and Future Work

- Empirical study of custom allocators
 - Lea allocator often as fast or *faster*
 - **Non-region custom allocation ineffective**
- Reap: region performance without drawbacks
- Future work:
 - Reduce space with per-page bitmaps
 - Combine with scalable general-purpose allocator (e.g., Hoard)



Software

<http://www.cs.umass.edu/~emery>

(Reap: part of Heap Layers distribution)

<http://g.oswego.edu>

(DLmalloc 2.7.0)



**If You Can Read This,
I Went Too Far**



Backup Slides



Experimental Methodology


- Comparing to general-purpose allocators
 - Same semantics: no problem
 - E.g., disable per-class allocators
 - Different semantics: use *emulator*
 - Uses general-purpose allocator
 - Adds bookkeeping to support region semantics



Why Did They Do That?


- Recommended practice
- Premature optimization
 - Microbenchmarks vs. actual performance
- Drift
 - Not bottleneck anymore
- Improved competition
 - Modern allocators are better

Reaps as Regions: Runtime



- Reap performance nearly matches regions

Using Reap as Regions



Reap performance nearly matches regions



Drawbacks of Regions

- Can't reclaim memory within regions
 - Bad for long-running computations, producer-consumer patterns, “malloc/free” programs
 - ⇒ **unbounded memory consumption**
- Current situation for Apache:
 - vulnerable to denial-of-service
 - limits runtime of connections
 - limits module programming



Use Custom Allocators?

- Strongly recommended by practitioners
- Little hard data on performance/space improvements
 - Only one previous study [Zorn 1992]
 - Focused on just one type of allocator
 - Custom allocators: waste of time
 - Small gains, bad allocators
- Different allocators better? Trade-offs?



Kinds of Custom Allocators

- Three basic types of custom allocators
 - Per-class
 - Fast
 - Custom patterns
 - Fast, but very special-purpose
 - Regions
 - Fast, possibly more space-efficient
 - Convenient
 - Variants: *nested*, *obstacks*

Optimization Opportunity

